

Technical Paper Proposal: Placement Algorithms for Heterogenous FPGAs

Brian B Cheng
Department of Electrical and Computer Engineering

1 Keywords

- FPGA, EDA, Synthesis, Placement, Routing, Parallel, Optimization

2 Proposal

A compiler takes a high-level program like a Java or Python file and assembles a corresponding machine code which is ready to be executed on a CPU. In similar fashion, an Electronic Design Automation tool (EDA) takes a high-level description of a digital system written in Verilog or VHDL and produces a bitstream which is ready to be deployed to an FPGA. In a superficial way, compilers and EDAs perform the same task for engineers but in different industries. One for software, the latter for hardware.

Software compilers have evolved to become highly optimized, significantly boosting the productivity of software developers. The compilers are so robust that developers rarely need to worry about the correctness or efficiency of their machine code, and are so refined that it all happens in a matter of milliseconds to minutes, even for large projects. The speed of compilation enables rapid debugging and implementation of design changes and allows developers to iterate through dozens or even hundreds of design cycles per day. Instead of tinkering around with assembly code, software engineers can comfortably focus on the high-level abstractions of their design.

The compilation stages mainly consists of lexing, parsing, and instruction optimization. While these stages can be complex from a design standpoint, they are, mathematically, trivial. The EDA stages are despairingly complex by comparison. First, the FPGA engineer creates and describes a digital design using a hardware description language (HDL) like Verilog or VHDL. Then the engineer submits the design entry to an EDA like Vivado or Quartus to perform three stages of automated design. These are: synthesis, placement, and routing. All three are NP-hard problems.

For even modest designs that utilize a small percentage of the available resources on the FPGA, the EDA will spend a minimum of a couple minutes

synthesizing and implementing the design. The EDA run time increases exponentially with the scale of the project. Projects that utilize 80% or more of the FPGA's resources may run for hours, and for the high-end devices like Xilinx's Kintex FPGAs which have millions of logical elements, can run for days. At high utilization, there is a possibility that the EDA cannot fit the design onto the device due to routing congestion even when it has enough logical elements to implement the design. The EDA might only reach this conclusion after attempting to fit the design for several hours.

One of the most common complaints from new FPGA engineers, especially those coming from software development, is that the FPGA toolchains are overall slow and heavy. This is because the problems that the EDA must solve are inherently complex. Unlike in software development where an error in the program can be patched and recompiled in a matter of minutes, EDA runtimes lasting several hours means that an engineer can only go through a couple design cycles per day. An FPGA engineer must be very precise with their coding and practice thorough verification of their design before submitting it to the EDA. This overall heightens the barrier of entry into FPGA development and contributes to a shortage of qualified FPGA engineers and a limitation of their productivity.

This paper aims to study one particular pillar of this barrier: the placement stage. First, it will review prerequisite knowledge for placement algorithms. This includes graph theory to represent logical elements of a digital design and the wired connections between them, and a corresponding optimization problem with a cost function which we seek to minimize.

In mathematics, a graph is modeled as $G = (V, E)$, where V represents the set of vertices and E represents the set of edges between them. A hypergraph is denoted as $G_H = (V_H, E_H)$, where V_H represents the set of vertices, and E_H represents the set of hyperedges, which are edges that can connect more than two vertices. Electronic circuits are represented as hypergraphs as because a voltage source can have one or many sinks. In FPGA design, the hypergraph is often transform into something called a "netlist", which is a hypergraph that has been flattened down into a graph. For every hyperedge in the hypergraph, transform the hyperedge into a collection of 2-pin edges - one source to one sink. There are different strategies for this reduction - unoriented star model, clique model, etc.. A netlist is a graph $G = (V, E)$, where V represents the components (e.g., logic blocks or gates), and E represents the connections or nets between them.

The vast majority of existing placers use the Half Perimeter Wire Length (HPWL) cost function or one of its various flavors - manhattan distance, euclidian distance, beta-regularized, etc.. The basic HPWL models the cost function as the summation of the manhattan lengths of all nets in the netlist.

This paper will primarily use Xilinx's open-source API, RapidWright to implement various custom placement algorithms. RapidWright is a Java API that provides backend access to Xilinx's EDA, Vivado.

$$\text{HPWL} = \sum_{i=1}^n \left(\max_{j \in \text{net}_i} (x_j) - \min_{j \in \text{net}_i} (x_j) + \max_{j \in \text{net}_i} (y_j) - \min_{j \in \text{net}_i} (y_j) \right) \quad (1)$$

Then, the paper will survey the current landscape of placement algorithms and strategies. This includes Simulated Annealing (SA) which is stochastic, but will focus on Analytical Placement (AP).

3 Abbreviations

- **FPGA**: Field Programmable Gate Array
- **VLSI**: Very Large Scale Integration
- **EDA**: Electronic Design Automation
- **VHSIC**: Very High Speed Integrated Circuits
- **HDL**: Hardware Description Language
- **VHDL**: VHSIC HDL
- **HLS**: High Level Synthesis: Generating synthesizable HDL from high-level software languages. A company might want to have a software engineer write C or C++ code and have a program translate it into synthesizable Verilog. This can boost productivity and save the company the need to hire a hardware engineer.
- **IP**: Intellectual Property: In FPGA context, this means pre-built modules or subsystems like a hardened microprocessor or Ethernet controller. These are usually proprietary.
- **SoC**: System on Chip: An FPGA device (chip) that features hardened IP in addition to the programmable logic fabric.
- **PL-PS**: Programmable Logic - Processing System: A design that utilizes the on-chip hard microprocessor in conjunction with the programmable logic fabric.
- **EDIF**: Electronic Design Interchange Format
- **HPWL**: Half Perimeter Wire Length

4 Ideas

- **FPGA**: Field Programmable Gate Array
 - FPGA Vendors:
 - AMD-Xilinx (~50% FPGA vendor market share)

- Intel-Altera (~35% share)
 - Lattice
 - Microsemi
- **EDA:** Electronic Design Automation
 - Proprietary software for FPGA and VLSI development:
 - Xilinx - Vivado (Design + Simulation) + Vitis (HLS + PL-PS code-sign)
 - Altera - Quartus (Design) + ModelSim (Simulation)
 - Synopsis (VLSI)
 - Cadence (VLSI)
 - Open source software for FPGA development:
 - **VTR:** Simulated Annealing placer for FPGAs. Popular among researchers who study placement techniques. Commonly referred to as an "academic placer".
 - **OSS-CAD:** a full-flow software suite that includes ABC synthesis, Yosys synthesis, Yosys nextpnr.
 - **AMF-Placer:** Analytical Placer for FPGAs
 - **RapidWright:** Semi-open source Java API that provides backend access to Xilinx Vivado EDA using design checkpoints.
 - **RapidLayout:** Hard Block Placer for Systolic Arrays. Built with RapidWright.
 - **RapidStream:** HLS Placer. Built with RapidWright.
 - **DREAMPlace:** GPU-powered deep learning placement for VLSI.
 - **DREAMPlaceFPGA:** DREAMPlace, adapted to FPGAs via the RapidWright API.
- **Synthesis**
 - Takes a design written in a high-level HDL like VHDL or Verilog and "synthesizes" a **logical netlist** out of it.
 - The logical netlist is usually generated as an EDIF, JSON, or a low-level Verilog file.
 - The netlist describes the necessary basic elements of logic (BELs) and the wired connections between them that are necessary to implement the design.
- **Placement**
 - Takes the **logical netlist** and produces a **physical netlist**.
 - For each BEL in the netlist, assign the BEL to a Cell, Site, and Tile on the physical FPGA device.
- **Routing**

- Takes the **physical netlist** and maps the connections between BELs onto wires, interconnects, and switchboxes on the FPGA.




Figure 1: Landscape of VLSI placement techniques (Spindler) [1]

Foundational Exploration		Modern Developments		Recent Progress
<1970s - 1980s	1980s - 1990s	1990s - 2010s		>2010s
Partitioning	Simulated Annealing	Min-Cut (Multi-level)	Analytic Techniques	
			Quadratic / Force-directed	Nonlinear Optimization
Breuer Dunlop and Kernighan Quadratic Assignment Cheng and Kuh PROUD † Cadence/QPlace*	TimberWolf/VPR † Dragon	FengShui Capo † Capo+Rooster	GORDIAN GORDIAN-L BonnPlace * mFar Kraftwerk † FastPlace3/RQL * Warp3	APlace2 Naylor/Synopsys * NTUPlace3 † mPL6 †
				Analytic Techniques Quadratic POLAR * SimPL/ComPLx MAPLE * Nonlinear ePlace
				† Used in industry * Commercial Placer

Figure 2: Historical timeline of VLSI placement techniques (Markov) [2]

This is a citation for AMFPlacer. [3]

References

- [1] P. Spindler, U. Schlichtmann, and F. M. Johannes, “Kraftwerk2—a fast force-directed quadratic placement approach using an accurate net model,” *Trans. Comp.-Aided Des. Integ. Cir. Sys.*, vol. 27, p. 1398–1411, aug 2008.
- [2] I. L. Markov, J. Hu, and M.-C. Kim, “Progress and challenges in vlsi placement research,” *Proceedings of the IEEE*, vol. 103, no. 11, pp. 1985–2003, 2015.
- [3] T. Liang, G. Chen, J. Zhao, S. Sinha, and W. Zhang, “Amf-placer: High-performance analytical mixed-size placer for fpga,” in *2021 IEEE/ACM International Conference on Computer-Aided Design (ICCAD)*, pp. 1–6, 2021.
- [4] A. Mishra, N. Rao, G. Gore, and X. Tang, “Architectural exploration of heterogeneous fpgas for performance enhancement of ml benchmarks,” in *2023 IEEE Asia Pacific Conference on Circuits and Systems (APCCAS)*, pp. 232–235, 2023.
- [5] U. Farooq, H. Parvez, H. Mehrez, and Z. Marrakchi, “Exploration of heterogeneous fpga architectures,” *Int. J. Reconfig. Comput.*, vol. 2011, jan 2011.
- [6] C. Lavin, M. Padilla, S. Ghosh, B. Nelson, B. Hutchings, and M. Wirthlin, “Using hard macros to reduce fpga compilation time,” in *2010 International Conference on Field Programmable Logic and Applications*, pp. 438–441, 2010.
- [7] C. Lavin and A. Kaviani, “Rapidwright: Enabling custom crafted implementations for fpgas,” in *2018 IEEE 26th Annual International Symposium on Field-Programmable Custom Computing Machines (FCCM)*, pp. 133–140, 2018.
- [8] Y. Zhou, P. Maidee, C. Lavin, A. Kaviani, and D. Stroobandt, “Rwroute: An open-source timing-driven router for commercial fpgas,” *ACM Trans. Reconfigurable Technol. Syst.*, vol. 15, nov 2021.
- [9] C. Lavin and E. Hung, “Invited paper: Rapidwright: Unleashing the full power of fpga technology with domain-specific tooling,” in *2023 IEEE/ACM International Conference on Computer Aided Design (ICCAD)*, pp. 1–7, 2023.
- [10] A. A. Kennings and I. L. Markov, “Analytical minimization of half-perimeter wirelength,” in *Proceedings of the 2000 Asia and South Pacific Design Automation Conference, ASP-DAC '00*, (New York, NY, USA), p. 179–184, Association for Computing Machinery, 2000.

- [11] S. Areibi, G. Grewal, D. Banerji, and P. Du, “Hierarchical fpga placement,” *Canadian Journal of Electrical and Computer Engineering*, vol. 32, no. 1, pp. 53–64, 2007.
- [12] M. Gort and J. H. Anderson, “Analytical placement for heterogeneous fpgas,” in *22nd International Conference on Field Programmable Logic and Applications (FPL)*, pp. 143–150, 2012.
- [13] J. Chen, W. Zhu, J. Yu, L. He, and Y.-W. Chang, “Analytical placement with 3d poisson’s equation and admm based optimization for large-scale 2.5d heterogeneous fpgas,” in *2019 IEEE/ACM International Conference on Computer-Aided Design (ICCAD)*, pp. 1–8, 2019.
- [14] S. Dhar, L. Singhal, M. Iyer, and D. Pan, “Fpga accelerated fpga placement,” in *2019 29th International Conference on Field Programmable Logic and Applications (FPL)*, pp. 404–410, 2019.
- [15] R. S. Rajarathnam, M. B. Alawieh, Z. Jiang, M. Iyer, and D. Z. Pan, “Dreamplacefpga: An open-source analytical placer for large scale heterogeneous fpgas using deep-learning toolkit,” in *2022 27th Asia and South Pacific Design Automation Conference (ASP-DAC)*, pp. 300–306, 2022.
- [16] H. Long, Y. Bai, Y. Li, J. Wang, and J. Lai, “Optimizing wirelength and delay of fpga tile through floorplanning based on simulated annealing algorithm,” in *2023 IEEE 15th International Conference on ASIC (ASICON)*, pp. 1–4, 2023.
- [17] T. Liang, G. Chen, J. Zhao, S. Sinha, and W. Zhang, “Amf-placer 2.0: Open source timing-driven analytical mixed-size placer for large-scale heterogeneous fpga,” *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, pp. 1–1, 2024.
- [18] B. Ray, A. R. Tripathy, P. Samal, M. Das, and P. Mallik, “Half-perimeter wirelength model for vlsi analytical placement,” in *2014 International Conference on Information Technology*, pp. 287–292, 2014.
- [19] B. N. B. Ray, S. K. Mohanty, D. Sethy, and R. B. Ray, “Hpwl formulation for analytical placement using gaussian error function,” in *2017 International Conference on Information Technology (ICIT)*, pp. 56–61, 2017.
- [20] Y. Lin, S. Dhar, W. Li, H. Ren, B. Khailany, and D. Z. Pan, “Dreamplace: Deep learning toolkit-enabled gpu acceleration for modern vlsi placement,” in *2019 56th ACM/IEEE Design Automation Conference (DAC)*, pp. 1–6, 2019.